



The Problem

In warehouse operations, picking each order individually is inefficient. Instead, we group compatible orders into **waves** so that their items can be collected together through shorter and more efficient routes. The goal is to decide which orders should form the next wave to maximize **picking productivity** — that is, to collect as many products as possible while visiting as few aisles as needed.

Let O be the set of pending orders, I_o the items requested in order $o \in O$, A the set of aisles, and $A_i \subseteq A$ the aisles containing item i . Each order o requests u_{oi} units of item i , and each aisle a holds u_{ai} units. Wave size is bounded by LB, UB lower and upper bounds on the total number of items in the wave.

We want to select:

$$O' \subseteq O \quad (\text{orders in the wave}), \quad A' \subseteq A \quad (\text{aisles to visit})$$

so as to maximize the ratio between collected units and visited aisles:

$$\max_{O', A'} \frac{\sum_{o \in O'} \sum_{i \in I_o} u_{oi}}{|A'|}$$

subject to:

$$LB \leq \sum_{o \in O'} \sum_{i \in I_o} u_{oi} \leq UB \quad (1)$$

$$\sum_{o \in O'} u_{oi} \leq \sum_{a \in A'} u_{ai}, \quad \forall i \in I_o, \quad o \in O' \quad (2)$$

A pair (O', A') satisfying (1)–(2) defines a feasible wave, and the optimal wave maximizes the productivity ratio above.

An Exact Parametric Approach

We apply a parametric algorithm based on Dinkelbach's method to solve the fractional objective. The problem is reformulated as finding the root of a convex function:

$$\phi(\lambda) = \max_{x \in X} \{f(x) - \lambda g(x)\},$$

where $f(x)$ and $g(x)$ are linear. Each evaluation of $\phi(\lambda)$ is obtained by solving a linear program (LP). In our context, $f(x)$ represents the total number of picked units, and $g(x)$ the number of aisles visited.

Dinkelbach's method can be interpreted as a Newton-type root-finding approach, since it follows the update rule:

$$\lambda_{k+1} = \lambda_k - \frac{\phi(\lambda_k)}{\phi'(\lambda_k)} = \lambda_k + \frac{\phi'(\lambda_k)}{g(x_k)},$$

where x_k is the optimal solution of the LP at iteration k . This iterative process continues until $\phi(\lambda_k) = 0$, ensuring convergence to the optimal productivity ratio λ^* [1].

Warm start

The Dinkelbach algorithm can benefit from a high quality initial solution, and thus we consider two simple greedy strategies to obtain them. The first one prioritizes picking aisles of a big *size* (i.e. those $a \in A$ that maximize $\sum_{i \in I_o} u_{ai}$) while the second one prioritizes aisles with high *diversity* (i.e. those $a \in A$ that maximize $|\{i \in I_o : u_{ai} > 0\}|$). As seen in Figure 1 optimal solutions have these type of aisles.




Figure 1. Aisles sorted by size and diversity and optimal aisles for instance 8 from dataset A.

Greedy vs. Optimal Solutions

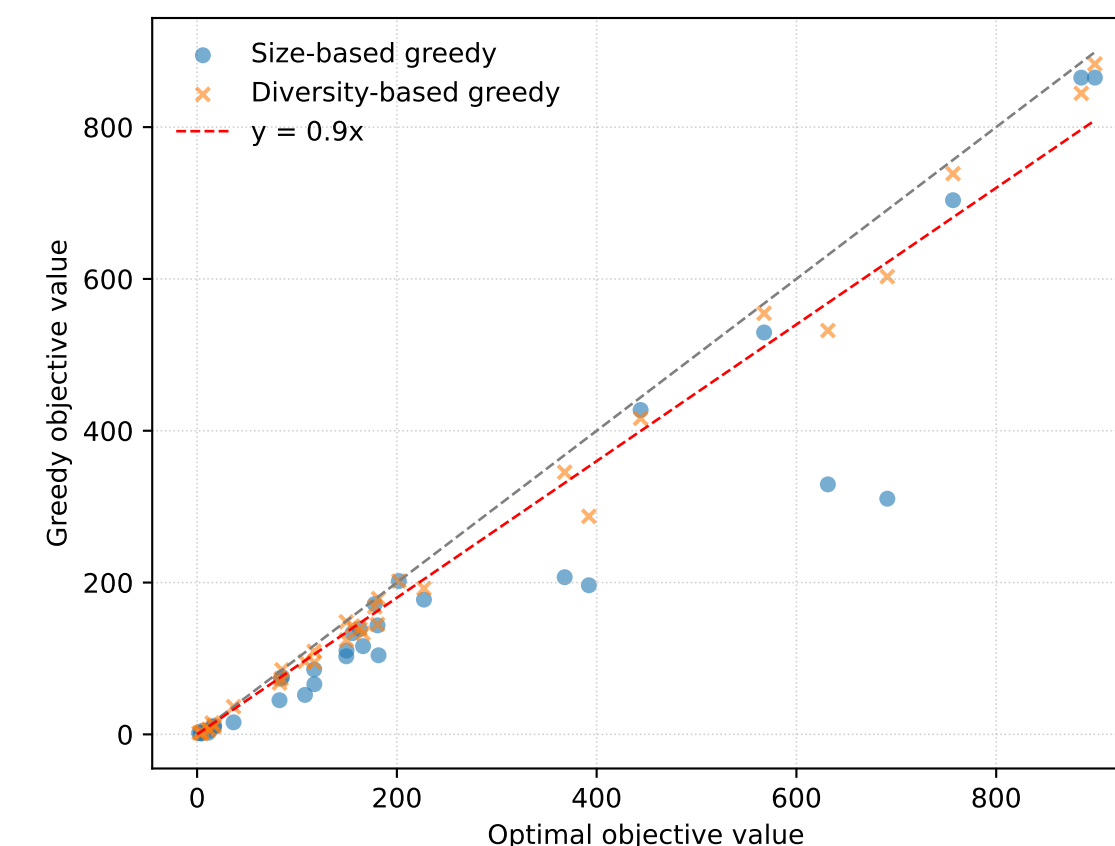


Figure 2. Comparison between the optimal values and the results given by the greedy algorithms.

Local Search Iteration

Experimentally we see that after the first iterations of the Dinkelbach algorithm the aisles do not change that much. This motivates the idea of enforcing the condition that the next solution found must differ from the current one by at most C aisles, which can be captured using the linear constraint

$$\sum_{a \in A_i} y_a - \sum_{a \in A \setminus A_i} y_a \leq C$$

where A_i denotes the aisles from the solution at iteration i and y_a is the indicator variable for aisle a . If an iteration includes this constraint it is faster at the potential cost of optimality. Thus, in our final algorithm we alternate between local and non-local iterations.

Implementation details

When implementing the parametric approach, we have to consider certain details to improve the efficiency of the algorithm. First of all, the set up of CPLEX's parameters, like the use of threads and RINS heuristic. On the other hand, we notice that keeping alive the structures like variables and constraints of the model through the iterations made a significant improvement in the time.

In addition, at the first iterations, we use a gap tolerance relatively high to achieve rapid results. However, as the algorithm comes near the optimal value, we lower down the gap tolerance until it reaches 0%. Although not only the gap tolerance is modified, we use an extra structure to maintain a time limit through the iterations, used to help prevent CPLEX from exploring useless branches and forcing it to start over from its current best when the limit is reached.

Plots

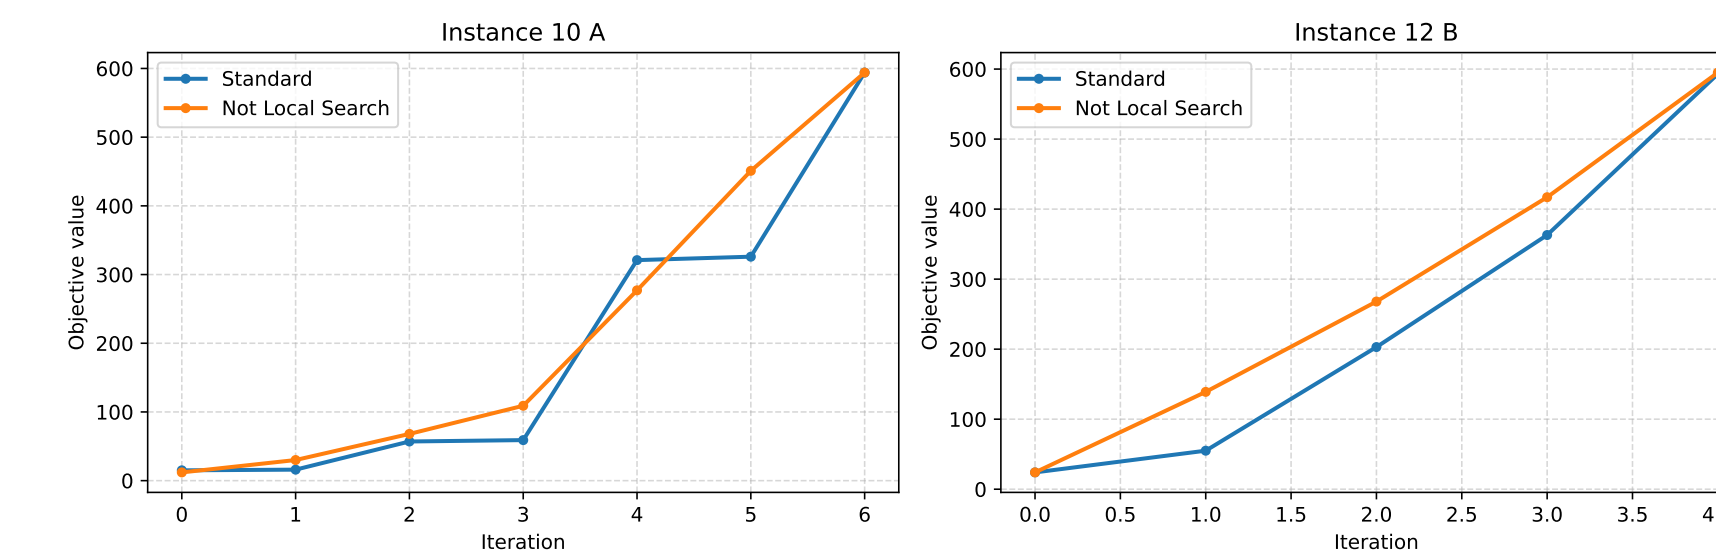


Figure 3. Comparison of the time through iterations using and not using the warmstart




Figure 4. Comparison of the objective value through iterations using and not using the warmstart

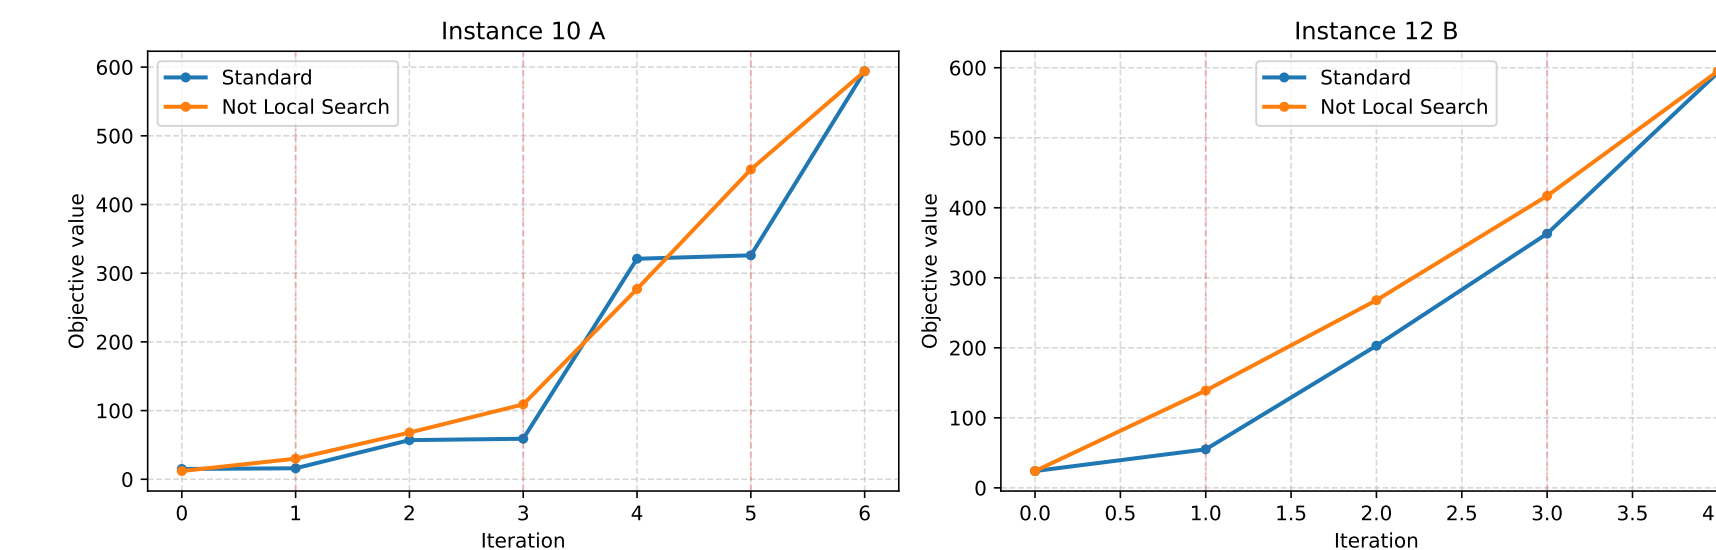


Figure 5. Comparison of the time through iterations using and not using the warmstart




Figure 6. Comparison of the objective value through iterations using and not using the warmstart